

Assignment: Haskell, Java Stream API & Python

Version 1.0 - December 15, 2020

Instructions

This assignment is made of three parts, consisting of exercises on Haskell, the Stream API of Java and Python, respectively.

Warning: This document is subject to changes. Check that you are reading always the most recent version.

Part 1 - Implementing multisets in Haskell

This assignment requires you to implement a type constructor providing the functionalities of [Multisets](#). Your implementation must be based on the following *concrete* Haskell definition of the `ListBag` type constructor:

```
data ListBag a = LB [(a, Int)]
    deriving (Show, Eq)
```

Therefore a `ListBag` contains a list of pairs whose first component is the actual element of the multiset, and the second component is its *multiplicity*, that is the number of occurrences of such element in the multiset. A `ListBag` is **well-formed** if it does not contain two pairs (v, k) and (v', k') with $v = v'$.

Exercise 1: Constructors and operations

The goal of this exercise is to write an implementation of multisets represented concretely as elements of the type constructor `ListBag`. As described below, the proposed implementation must be well documented and must pass the provided tests.

- Using the type constructor `ListBag` described above, implement the predicate `wf` that applied to a `ListBag` returns `True` if and only if the argument is well-formed. Check that the inferred type is `wf :: Eq a => ListBag a -> Bool`.

Important: All the operations of the present exercise that return a `ListBag` bag must ensure that the result is well-formed, i.e., that `wf bag == True`.

- Implement the following constructors:
 - `empty`, that returns an empty `ListBag`
 - `singleton v`, returning a `ListBag` containing just one occurrence of element `v`
 - `fromList lst`, returning a `ListBag` containing all and only the elements of `lst`, each with the right multiplicity
- Implement the following operations:
 - `isEmpty bag`, returning `True` if and only if bag is empty
 - `mul v bag`, returning the multiplicity of `v` in the `ListBag` `bag` if `v` is an element of `bag`, and `0` otherwise
 - `toList bag`, that returns a list containing all the elements of the `ListBag` `bag`, each one repeated a number of times equal to its multiplicity
 - `sumBag bag bag'`, returning the `ListBag` obtained by adding all the elements of `bag'` to `bag`

Testing: The attached files [testEx1.hs](#) and [testEx12.hs](#) contain some tests that can be used to check the correctness of the implemented functions. To run the tests, it is probably necessary to install the **Test.HUnit** Haskell module. Next it is sufficient to load file `testEx1.hs` or `testEx12.hs` in the interpreter, and execute `main`.

Note that solutions that do not pass such tests will not be evaluated, and a revision will be requested.

Solution format: A Haskell source file called `Ex1.hs` containing a [Module \(see Section "Making our own modules"\)](#) called `Ex1`, defining the data type `ListBag` (copy it from above) and *at least* all the functions described above. The module can include other functions as well, if convenient. **Note:** The file has to be adequately commented, and each function definition must be preceded by its type, as inferred by the Haskell compiler.

Exercise 2: Mapping and folding

The goal of this exercise is to experiment with class constructors by adding some functions to the module developed for [Exercise 1](#).

1. Define an instance of the constructor class `Foldable` for the constructor `ListBag` defined in [Exercise 1](#). To this aim, choose a minimal set of functions to be implemented, as described in the documentation of `Foldable`. Intuitively, folding a `ListBag` with a binary function should apply the function to the elements of the multiset, ignoring the multiplicities.
2. Define a function `mapLB` that takes a function `f :: a -> b` and a `ListBag` of type `a` as an argument, and returns the `ListBag` of type `b` obtained by applying `f` to all the elements of its second argument.
3. Explain (in a comment in the same file) why it is not possible to define an instance of `Functor` for `ListBag` by providing `mapLB` as the implementation of `fmap`.

Solution format: A Haskell source file `Ex2.hs` containing a module called `Ex2`, which imports module `Ex1` and includes **only** the new functions defined for this exercise. **Note:** The file has to be adequately commented, and each function definition has to be preceded by its type, as inferred by the Haskell compiler.

Part 2 - A Map-Reduce framework exploiting the Java Stream API

The Map-Reduce paradigm is widely used for processing huge amounts of data in a parallel and distributed setting. In this assignment, students are required to implement a simple software framework providing the functionalities of Map-Reduce, but ignoring the aspects of parallelism and distribution. As a proof of concept, two simple working instances of the framework should be implemented as well.

For an introduction to the Map-Reduce framework see the paper [MapReduce: Simplified Data Processing on Large Clusters](#). For a presentation of Map-Reduce as Software Framework, identifying the *hot spots*, see <https://en.wikipedia.org/wiki/MapReduce#Dataflow> (since we ignore the distribution aspects, you can ignore the *Partition function*.)

Solution format: An archive `MapReduce-<yourSurname>.zip` containing the Java files implementing Exercises 3, 4, and (optionally) 5. If you use NetBeans, please send in the archive the entire project.

Exercise 3 - The framework

Following the guidelines presented in the lesson of October 23, 2020 (see <http://pages.di.unipi.it/corradini/Didattica/AP-20/index.html#framework>), and more specifically the *Template Method design pattern*, implement in Java a Map-Reduce software framework providing the functionalities described in the above documentation and respecting the following constraints:

1. For key/value pairs, the framework must use the attached class [Pair.java](#) (you can change its package, but nothing else).

2. The hot spots of the framework are the methods `read`, `map`, `compare`, `reduce` and `write`.
3. The framework must use, when possible, the Stream API. For example, `map` takes a stream of key-value pairs as argument and returns a stream of key-value pairs (types of argument and result may differ, of course).

Exercise 4 - Counting words

By instantiating the framework, implement a program that counts the occurrences of words of length greater than 3 in a given set of documents, respecting the following constraints:

1. The program should ask the user for the absolute path of the directory where documents are stored. Only files ending in `.txt` should be considered.
2. The `read` function must return a stream of pairs (`fileName`, `contents`), where `fileName` is the name of the text file and `contents` is a list of strings, one for each line of the file. For the `read` function you can exploit the enclosed class [Reader.java](#) in the way you prefer.
3. The `map` function must take as input the output of `read` and must return a stream of pairs containing, for each word (of length greater than 3) in a line, the pair (`w`, `k`) where `k` is the number of occurrences of `w` in that line.
4. The `compare` function should compare strings according to the standard alphanumeric ordering. (The result should adhere to the standard Java conventions, see the `compareTo` method of interface [Comparable](#).)
5. The `reduce` function takes as input a stream of pairs (`w`, `lst`) where `w` is a string and `lst` is a list of integers. It returns a corresponding stream of pairs (`w`, `sum`) where `sum` is the sum of the integers in `lst`.
6. The `write` function takes as input the output of `reduce` and writes the stream in a CSV (Comma Separated Value) file, one pair per line, in alphanumeric ordering. For the `write` function you can exploit the enclosed class [Writer.java](#) in the way you prefer.

For testing the program you can use the enclosed archive [Books.zip](#) which contains parts of some famous books as downloaded from the pages of the [Gutenberg Project](#). (before the site became inaccessible from Italy, see [Raffaele Angius, Perché il Progetto Gutenberg sarà sotto sequestro per sempre](#)).

Exercise 5 - [Optional] Producing an Inverted Index

By instantiating the framework, implement a program that generates an *Inverted Index* (for words of length greater than 3). That is, given as input the absolute path of a directory, the program prints in a CSV file for each word `w` (of length greater than 3) appearing in the `.txt` documents of the directory, a line `w, filename, line` if `w` appears in line number `line` of file `filename`. The lines should be sorted in the natural way.

Part 3 - Benchmarking functions in Python, with Multithreading

You want now to benchmark your Python functions and you want to make sure that it is done in the right way (hopefully... :P)

Solution format: A single Python file called `benchmark.py` containing the solutions to the following three exercises.

Exercise 6 - A decorator for benchmarking

Define a Python decorator called `benchmark`. When invoking a function `fun` decorated by `benchmark`, `fun` is executed possibly several times (discarding the results) and a small table is printed on the standard output including the average time of execution and the variance.

The exact behaviour of the `benchmark` decorator is ruled by the following optional parameters:

- `warmups`: The number of warm-up invocations to `fun` (i.e. invocations whose timing must be ignored) (default: `warmups = 0`);
- `iter`: The number of times `fun` must be invoked and whose timing must be taken into account for the final metrics (default: `iter = 1`);
- `verbose`: Whether the execution should be verbose (i.e. if it must print the timing of each warm-up round and invocation) or not (default `verbose = False`);
- `csv_file`: A CSV file name where the benchmark information must be written. The header of the file will be in the form `run num, is warmup, timing` with the intuitive meaning (default `csv_file = None`, meaning that the benchmark information is only displayed on screen).

Exercise 7 - Testing the decorator with multithreading

Test your implementation by also evaluating the effectiveness of multithreading in Python.

Using the [threading](#) module of the Standard Library and exploiting the benchmark decorator, write a function `test` that executes a function `f` (passed as parameter) with varying numbers of iterations and degrees of parallelism. More precisely, the `test` first runs `f` 16 times on a single thread, then 8 times on two threads, then 4 times on 4 threads, and finally 2 times on 8 threads. The program must write the benchmarking information for the four scenarios in corresponding files named `f_<numthreads>_<numiterations>`.

Run the program using a function that computes the n -th Fibonacci number in the standard, inefficient, double recursive way (choose n carefully) .

Discuss briefly the results in a comment in the Python file.

Author: Andrea Corradini & Laura Bussi

Created: 2020-12-15 Tue 01:40

[Validate](#)